



AGENTIC SOFTWARE ENGINEERING FOR LEADERS

AI-Powered IDE Coding Agents for Organizations:
A Handbook of Best Practices

LUCIANO AYRES

Agentic Software Engineering for Leaders

AI-Powered IDE Coding Agents for Organizations: A Handbook of Best Practices

By [Luciano Ayres](#)

First Edition | Published March 23, 2025

Version 1.0

Table of Contents

Introduction.....	3
Agentic Software Engineering.....	3
What Is Agentic Software Engineering?.....	3
Why Is Agentic Software Engineering Relevant Today?.....	4
AI Is Transforming Software Development Workflows.....	5
Next-Generation AI Tools: From Autocomplete to Autonomous Agents.....	6
The Autonomous AI Pair Programmer.....	6
Deep Context Awareness.....	7
Agentic Assistance & Developer Flow.....	7
Introducing AI Tools to Your Team: Strategy and Change Management.....	9
Prepare the Ground with Education and Buy-In.....	9
Onboarding and Adoption Tips.....	10
Measuring the Impact of AI Adoption.....	11
Fostering a Culture of AI Literacy and Experimentation.....	13
Balancing Automation with Human Judgment.....	15
Leading by Example: How Managers and CTOs Can Leverage AI.....	17
Conclusion.....	19
About the Author.....	20
References.....	21

Introduction

Artificial Intelligence is rapidly reshaping how software is built. Modern development teams are experiencing a sea change in workflows, productivity, and even job satisfaction thanks to AI-powered tools. As a software engineering manager, aspiring leader, or CTO, it's imperative to understand how to strategically leverage these AI tools to boost your team's productivity, improve code quality, and enhance developer quality of life.

This handbook provides a roadmap: from understanding the AI transformation in coding, to introducing cutting-edge autonomous IDE agents like Cursor and Windsurf, to practical tips on change management and fostering an AI-ready culture. The goal is clear, strategic guidance to help you lead your team into this new era of software development.

Agentic Software Engineering

What Is Agentic Software Engineering?

Agentic Software Engineering represents a shift from traditional, reactive coding tools to autonomous, proactive systems. In this context, "agentic" refers to tools that can take independent actions on your behalf. These tools are not just passive assistants; they actively participate in the software development process by:

- **Initiating Tasks:** AI agents can autonomously tackle multi-step tasks like refactoring code, creating test suites, or even synthesizing new features from high-level requirements.
- **Operating in Shadow Workspaces:** For example, Cursor uses "shadow workspaces" to safely experiment with code changes in the background before proposing them to developers.
- **Deep Integration and Autonomy:** Tools like Windsurf's Cascade are designed to traverse multiple files and even interact with build systems, test frameworks, and version control, while keeping the developer in the loop for approval.

Why Is Agentic Software Engineering Relevant Today?

The relevance of agentic approaches in modern software engineering is multi-fold:

- **Increased Complexity:** Modern codebases are vast and interconnected. Agentic systems can analyze and modify code at scale, reducing the manual overhead for developers.
- **Faster Innovation Cycles:** With AI agents handling routine work, developers can focus on innovative solutions, driving faster product iteration and improvement.
- **Quality and Consistency:** Autonomous agents ensure that coding standards and best practices are consistently applied across the codebase. This leads to fewer bugs and more maintainable software.
- **Enhanced Developer Experience:** By offloading repetitive and mundane tasks, agentic systems help maintain developer flow and reduce burnout. Developers can dedicate more time to creative problem solving and high-impact work.
- **Competitive Advantage:** Organizations that embrace agentic software engineering are better positioned to deliver high-quality software faster, providing a critical edge in today's fast-paced tech environment.

In summary, Agentic Software Engineering is not just a buzzword—it's a practical, strategic approach that leverages the latest in AI to meet the challenges of modern development head-on. As you move towards integrating autonomous IDEs, keeping agentic principles in mind ensures you balance automation with human oversight, leading to superior outcomes.

AI Is Transforming Software Development Workflows

Not long ago, AI in coding sounded futuristic. Today, it's mainstream. In fact, a recent survey by GitHub found 92% of developers are now incorporating AI coding tools into their workflow, and 70% report significant benefits from using these tools. What does this mean for development workflows? In practice, AI assistants are accelerating numerous aspects of the software lifecycle:

- **Faster coding with fewer errors:** AI pair programmers (like GitHub Copilot and others) can autocomplete boilerplate code, suggest entire functions, and flag mistakes in real time. Studies suggest developers see a 20–30% improvement in development speed with code assistants. At the same time, AI's ability to catch common errors and offer fixes can lead to a 10–15% reduction in bugs reaching production. This means faster feature delivery and less time firefighting defects.
- **Automating the tedious tasks:** Repetitive tasks like writing unit tests or documentation are now easier to handle. AI can generate unit test templates and even infer edge-case scenarios, yielding 10–20% better test coverage on average. It can also produce documentation drafts from code, keeping docs up-to-date without consuming developer hours. By offloading grunt work to AI, developers can focus on higher-value design and problem-solving tasks.
- **Streamlined codebase understanding:** Large Language Models (LLMs) excel at searching and summarizing information. Developers can ask an AI chatbot integrated in their editor questions about the codebase ("Where is the login logic defined?") and get instant answers or navigations to relevant code. This reduces context-switching and speeds up onboarding on complex projects.
- **Real-time code review and quality checks:** AI agents can act as tireless reviewers, scanning new code for potential issues, style deviations, or security vulnerabilities. For example, AI-powered code review tools can highlight risky code paths or suggest refactoring opportunities as code is written. This continuous feedback loop helps maintain higher code quality standards from the outset.

In short, AI is becoming a co-developer in the team, handling rote work and augmenting human capabilities. Development workflows are shifting from a manual, labor-intensive process to one of human oversight and creative direction, with AI handling many execution details. As an engineering leader, recognizing this shift is crucial – it's not hype, but a fundamental change in how software is built. Those who embrace it can achieve faster cycle times, more reliable code, and happier developers; those who ignore it risk falling behind. As the next sections will show, the key is leveraging these tools strategically and responsibly.

Next-Generation AI Tools: From Autocomplete to Autonomous Agents

Early AI coding tools (like basic autocomplete or standalone code generators) were helpful, but today's AI development environments take it a step further. Tools such as Cursor and Windsurf represent a new wave of AI-integrated IDEs that are autonomous, context-aware, and agentic – meaning they don't just assist with coding; they can take actions on your behalf in a controlled way. Let's explore how these advanced tools are changing the developer experience:

The Autonomous AI Pair Programmer

Traditional code assistants react to prompts or keystrokes, but autonomous agentic IDEs can proactively complete end-to-end tasks. For example, Cursor's "Agent" mode is designed to carry out multi-step coding tasks with minimal intervention. As the company describes, "Cursor's agent mode can complete tasks end to end... while keeping programmers in the loop." In practice, this means you could ask Cursor's agent to implement a new feature or refactor a module, and it will generate the necessary code across files, step through execution, fix simple errors, and present the changes for your review. The developer stays informed and approves the final edits, but the heavy lifting of writing and iterating on code is handled by the AI.

Windsurf's Cascade agent autonomously tackling a coding task across multiple files, then presenting diffs for review. In this example, Cascade analyzed the codebase, made multi-file edits (with lines added/removed), and is ready to open diffs for the developer to inspect.

Windsurf's Cascade is another example of autonomous capability. It introduces an AI "flow" that can act like a senior engineer executing a task. In Cascade's "Write Mode," it functions much like an AutoGPT specialized for coding – it will create or modify multiple files, run scripts or tests, observe the results, and debug as needed, iterating until the task is complete. Crucially, these agent actions happen with the developer's oversight: Cascade asks for your approval before running commands or making persistent changes. This design balances autonomy with control, ensuring the AI doesn't run amok in your repo without you knowing.

Cursor has tackled the autonomy challenge with a concept called "Shadow Workspaces." Running an AI agent that can freely modify code can be dangerous if it directly alters your project (imagine an AI refactoring half your codebase incorrectly). Cursor's solution is to let the AI work in a background clone of your workspace. "To actually be helpful, the AI iteration needs to happen in the background, without affecting your coding experience. To achieve this, we implemented what we call the shadow workspace in Cursor."

In a shadow workspace, the AI can compile, run, and rewrite code, see the effects (like linter errors or test failures), and iterate multiple times – all invisible to the developer until the AI is confident in a solution. Only then are the proposed changes presented to the user for review and application. This clever approach lets the AI be truly agentic (it can “think” by trial-and-error) while preserving the developer’s flow and preventing any chaos in the main workspace.

Deep Context Awareness

Another hallmark of these advanced tools is contextual awareness of the entire codebase and beyond. Unlike simpler autocomplete that only sees the current file or a few recent lines, Cursor and Windsurf index your whole project and even external resources. They leverage techniques like semantic search and embeddings to “read” your repository. The payoff is that the AI can make globally informed suggestions – for instance, renaming a function used across dozens of files, or finding the right API call by searching usage examples in the codebase. Windsurf’s documentation highlights “deep contextual awareness [that] allows you to run Cascade on production codebases and still get relevant suggestions.” In practice, the AI is less likely to produce nonsensical or out-of-scope code because it understands the larger context of your software.

These tools also integrate external knowledge sources. Both Cursor and Windsurf enable “mentions” of documentation or the web in AI prompts. For example, a developer can pull in an official library doc page or even allow a web search if the AI needs more information. “Cascade can intuitively parse and analyze web pages and documentation in real time, providing relevant and actionable context for your code.” This context merging means the AI agent can answer questions like “How do we use this new framework’s API here?” by consulting docs, or it can migrate code to a new library by following an online example it just read – all within the IDE. The result is fewer interrupts for the developer (no more manual Googling for that error message; the AI might do it for you) and a more seamless coding session.

Agentic Assistance & Developer Flow

Importantly, these AI-augmented IDEs still keep the developer in charge while providing ever-present assistance. They act as intelligent copilots that can take initiative when asked. The Windsurf team describes their approach as combining “collaborate with you like a Copilot and tackle complex tasks independently like an Agent. The AI is completely in sync with you, every step of the way.” This means the AI is aware of your current focus (cursor position, open files, recent actions) and can proactively help.

For instance, Windsurf’s “Supercomplete” feature analyzes what your next action might be – beyond just completing the current line of code – and prepares to assist accordingly. If you define a new function, the AI might anticipate that you’ll need to call it somewhere or write tests for it, and it can suggest those next steps. Similarly, Cursor’s predictive features attempt to guess the next

edit you might make and pre-suggest it, almost like an IDE that reads your mind (users have described moments where “Cursor is steps ahead of my brain, proposing multi-line edits so I type ‘tab’ more than anything else.”).

Beyond code suggestions, agentic IDEs can execute commands and handle tooling on behalf of developers. For example, if you tell the AI, “Generate and run a database migration for adding a new column,” a tool like Windsurf Cascade can write the migration script and execute the relevant CLI commands to apply it (in a safe, user-approved manner). These AI agents can thus interface with your build system, tests, linters, and version control. In fact, Windsurf will even auto-fix issues it encounters: “If Cascade generates code that doesn’t pass a linter, then Cascade will automatically fix the errors.” This kind of end-to-end handling reduces back-and-forth for developers. Instead of manually running a linter, seeing errors, fixing them, and re-running, the AI closes the loop for you. The developer simply sees the final result: code that meets the team’s standards.

All of these capabilities contribute to a smoother developer experience. Programmers can remain “in flow” longer, with fewer distractions and manual micro-tasks, because the AI agent is taking care of them in the background. It’s as if each developer has a dedicated junior engineer who is extremely fast, always available, and can execute instructions across the whole project. The senior developer (your human engineer) can then focus on guiding the AI, making design decisions, and reviewing the AI’s output for correctness and quality. This partnership can be transformative for productivity – but to reap the benefits, engineering leaders need to introduce and manage these tools thoughtfully. In the next section, we’ll look at how to do exactly that.

Introducing AI Tools to Your Team: Strategy and Change Management

Even the most powerful AI tool won't help if your team doesn't adopt it or use it effectively. As a manager or tech leader, you play a key role in championing these tools, guiding your team through the change, and ensuring positive outcomes. Here's how to introduce AI coding assistants like Cursor or Windsurf to your engineering team strategically:

Prepare the Ground with Education and Buy-In

1. **Demystify AI for your team:** Start by setting the right context. Clearly communicate why you're bringing in AI assistance – for example, to free them from drudge work, to speed up release cycles, and to help catch errors early. Emphasize that AI tools are there to empower developers, not replace them. This helps alleviate common fears (“Will AI make me obsolete?”) and frames the tools as a positive augmentation. It's often worth sharing data or case studies: for instance, how other teams achieved faster development with fewer bugs after adopting AI helpers. Developers are naturally curious; show them what's in it for them (less boilerplate, more time for interesting problems) and you'll spark enthusiasm.
2. **Invest in training and AI literacy:** Don't just throw the tool at the team and hope for the best. Provide hands-on training sessions or tutorials on how to use the AI assistant effectively. Cover basics like prompting techniques, the tool's features, and its limitations. Encourage developers to understand why the AI suggests certain code – fostering understanding is crucial so they can use suggestions judiciously. You might run an interactive workshop where the team walks through using Cursor's chat to solve a coding problem, or demonstrate how Windsurf's Cascade can automate a refactor. Pair new users with early adopters or “AI champions” on the team who can mentor and answer questions in the first few weeks. The goal is to build confidence and competence with the tool, turning initial novelty into everyday productivity.
3. **Start with a pilot project:** It's wise to introduce AI assistance gradually. Identify a low-risk pilot project or a subset of the team to trial the tool first. For example, pick a small feature development or an internal tool rewrite as the sandbox for using Cursor or Windsurf extensively. This pilot group can experiment, figure out what works well and what processes need adjustment, and then share their experiences with the rest of the team. By starting small, you create a safe learning environment and can work out any kinks in integration before scaling up. Collect feedback from the pilot – were there significant time savings? What challenges did developers encounter? Use those insights to refine how you'll roll it out more broadly.

4. Set clear guidelines and expectations: As you roll out the AI tool team-wide, establish best practices. For instance, define that all AI-generated code must be reviewed by a human before merging – this underscores that human judgment remains the final word (more on that later). You might create a short “AI Assistant Usage Guidelines” document covering things like: appropriate use cases (e.g. drafting code, writing tests, exploring solutions), when not to use the AI (e.g. for security-critical code without careful review), how to handle any sensitive data (ensuring the tool’s privacy settings are used), and coding style adherence. Setting these expectations early helps integrate the tool into your existing workflow.

Developers should view the AI as a powerful helper within known guardrails, rather than a mysterious black box. Encourage the team to treat AI suggestions as they would a colleague’s code review comments – valued input, but subject to scrutiny.

5. Lead by example: Adoption starts at the top. If you’re an engineering manager or tech lead who still occasionally writes code or reviews code, make a point of using the AI tool yourself and letting the team see that. Even if your coding contributions are minimal, you can use the AI to prototype solutions or analyze the codebase in team discussions. For example, during a planning meeting you might say, “I asked our AI assistant to outline a possible module design, and here’s what it came up with – let’s critique it.” This signals to the team that using the AI is encouraged and respected. It also gives you firsthand experience, so you understand its strengths and weaknesses and can better guide your team. Managers who embrace new tools openly tend to inspire their teams to follow suit.

Onboarding and Adoption Tips

Bringing AI into established workflows is a change management exercise. Here are a few additional tips to smooth the transition:

- Integrate into daily workflow: Encourage the team to use the AI tool in regular tasks (not just special occasions). This might mean installing the Cursor or Windsurf IDE for daily development, or integrating an AI assistant plugin into your primary source editor. The less context-switching required, the more likely devs will use it. Some teams start by using AI in code review or bug fix sessions, where the stakes of trying a suggestion are low. Over time, it becomes second nature to hit that AI auto-complete or ask the chatbot a question whenever needed.
- Address resistance with empathy: Not everyone will jump on the bandwagon immediately. Experienced developers might be skeptical of AI suggestions or worry it encourages bad habits. Listen to their concerns. Often, running a time-boxed trial helps – e.g., “Try using the AI for two days for even trivial tasks, and then let’s discuss what you liked or disliked.” Many skeptics become fans when they see it save them from tedious work (like writing getters and setters or boilerplate tests). Also reassure that using AI doesn’t diminish their

creativity or value; instead, it amplifies their impact by handling grunt work. Keep the dialogue open and non-judgmental.

- **Celebrate quick wins:** As the team begins using the AI tool, highlight any success stories. Did the AI help cut down a feature implementation from 5 days to 3? Did someone catch a tricky bug because the AI pointed it out? Share these anecdotes in team meetings or Slack channels. Early victories build momentum. You might even have team members do a short show-and-tell of a cool thing they did with Cursor or Windsurf. This positive reinforcement encourages others to experiment and builds confidence that the investment in AI is paying off.
- **Provide ongoing support:** The first few weeks of adoption are critical. Make sure support is available – whether it’s an internal guru on the team or resources from the tool’s vendor (documentation, office hours, etc.). Create a space for questions and knowledge sharing, like an internal wiki page or chat thread for AI tips and tricks. The more your engineers help each other (e.g., sharing useful prompt techniques or caveats they discovered), the faster everyone climbs the learning curve. As a manager, periodically check in on how folks are feeling about the tool and gather suggestions for improvement. This shows that you are attentive to their experience and open to adjusting course if needed.

Measuring the Impact of AI Adoption

To ensure these tools are delivering real value (and to make a case for continued investment), you’ll want to measure their impact on your team. Keep in mind that improvement will grow over time as the team becomes more adept with the AI. Here are some ways to gauge the effects:

- **Productivity metrics:** Look at your team’s velocity or throughput before vs. after AI adoption. Are user stories being completed faster? You might measure the average cycle time for a certain category of tasks (for example, how long does it take to implement a standard CRUD feature now versus before). If you track story points, has the team’s completed points trended upward? Be careful to account for the learning period; the biggest gains might come after a couple months of use. Some organizations have found dramatic improvements – developers can be up to 45% more productive with AI coding assistants, according to independent studies – but your mileage may vary. Focus on trends and improvements in your own context rather than absolute numbers.
- **Code quality indicators:** Measure whether code quality is improving alongside speed. For instance, track the number of defects found in QA or after release. Has it decreased now that AI is catching issues earlier? Monitor code review comments – are reviewers finding fewer basic mistakes or omissions? You could also use static analysis tools to see if common issues (null pointer risks, style violations, etc.) are reduced. The ideal outcome is faster output with equal or better quality. If you find quality slipping (maybe due to

over-reliance on the AI's first suggestion), that's a signal to add training or adjust processes (like enforcing thorough code reviews of AI-written code).

- **Developer satisfaction and engagement:** Don't forget the human side. Gather feedback on how the developers feel about using the AI tools. This can be informal discussions or a quick survey asking questions like "On a scale of 1-5, how much does the AI assistant help you in your daily work?" and "Do you feel it has improved your coding experience?" Improved morale, less frustration with boring tasks, and more time for creative work are all signs of success. Remember that developer happiness is a valid metric – a tool that makes engineers' lives easier can increase job satisfaction and reduce burnout. You might notice qualitative improvements: developers may say things like "I feel I can't live without this now" or "I get unstuck from problems faster than before." Those are big wins, even if they're hard to quantify. Conversely, if anyone feels the tool is hindering them, dig into why – maybe they need more training or perhaps a particular task isn't well-suited for AI.
- **Team and business outcomes:** At a higher level, assess if the team is delivering more value to end-users. Are you able to ship more features per quarter? Has customer-reported bug volume dropped? For CTOs, these broader metrics matter to justify the adoption of AI tools. Any increase in delivered customer value or maintenance of output with a smaller team (if that's a context) is a strong indicator that AI augmentation is working. When communicating upward to executives, tie the impact to business results (e.g., "Using the AI coding assistant, our team delivered the last release 2 weeks faster, which enabled the company to capture revenue from a new feature ahead of schedule").

Finally, be patient and give the team time to adapt. As GitLab's engineering productivity research notes, you can't measure AI's impact in just a week or two – teams need time to find their rhythm with AI assistants. Set evaluation checkpoints perhaps a month and a quarter into adoption to see how things are trending. And be ready to iterate: if a metric isn't moving in the desired direction, engage the team to understand why and adjust your approach. Adopting AI in development is a journey, and continuous improvement is part of it.

Fostering a Culture of AI Literacy and Experimentation

To truly reap the benefits of AI, it's not enough to install a tool – you need to embed AI into the DNA of your engineering culture. This means cultivating AI literacy, encouraging experimentation, and making continuous learning a norm. Here's how engineering leaders can create an AI-forward culture:

1. Build AI literacy across the team. Every team member, from junior developers to senior architects, should have a basic understanding of how these AI tools work and their potential. This doesn't require everyone to be a machine learning expert, but they should grasp concepts like what an LLM is, why prompting matters, and what the tool's strengths/weaknesses are. You might organize lunch-and-learn sessions on AI topics or share resources (articles, courses) for self-study.

One effective approach is to provide comprehensive training programs so everyone has a foundational understanding of AI principles and tools – this not only builds skills but also empowers engineers to spot new opportunities to use AI in the product. In practice, some companies have even created internal AI Academy programs or certification tracks to ensure a baseline of AI knowledge. The more fluent your team is in AI, the more effectively they'll wield these tools.

2. Encourage experimentation and play. To get past the hype and truly integrate AI, teams need to tinker and try things. Create a safe space for this experimentation. For instance, hackathons or “AI Fridays” can be great ways for engineers to play with AI features not directly related to sprint work. One engineering leader shared that they set aside Friday afternoons for anyone to explore generative AI, then “check in” on Slack to share what they attempted and learned. In a short time, small groups naturally formed to tackle fun mini-projects with AI, and they exchanged findings with excitement. You can even introduce friendly competitions – e.g., who can get the AI to solve a tricky coding challenge with the fewest prompts, or who can come up with the most useful prompt for generating test cases (with a small prize to spice it up). By creating a sandbox for AI, you signal that trying and even failing with these tools is not just allowed, but celebrated. This builds confidence and often surfaces creative use cases that management might not have thought of.

3. Allocate time for learning. It's hard for engineers to become AI-proficient if they're expected to figure it out only “on the side” of their regular duties. Recognize that learning AI-assisted development is an investment. Consider allocating a modest portion of work hours for AI learning – whether through formal training or self-directed practice. This could be as simple as letting each sprint include a half-day for developers to improve their AI usage skills or investigate new AI features/plugins. Some teams adopt the 90/10 rule: 90% of time on core tasks, 10% on innovation and skill growth (which AI falls under).

Leadership support is key: when managers explicitly allow and encourage taking time to master the new tools, it legitimizes the effort. It also accelerates the team's journey up the learning curve, leading to better productivity sooner. Remember, "experiential learning is absolutely essential for building AI literacy" – people learn by doing, so give them the time and space to do so.

4. Make knowledge sharing a norm. As your team experiments, create channels to share insights. Maybe an internal newsletter or weekly email where folks can post a cool trick they discovered ("Tip: Using @Codebase context in Cursor helped me quickly find all uses of function X"). Or dedicate 5 minutes in your team meeting for an "AI pro-tip of the week." This cross-pollination ensures everyone benefits from individual learnings and avoids duplicate blind alleys. You might end up building an internal knowledge base or FAQ for your AI tools as a living document. Over time, this collective wisdom significantly raises the whole team's capability. It also reinforces a culture where continuous improvement and learning are valued, which is healthy beyond just AI adoption.

5. Stay ethical and considerate. An AI-literate culture should also be one that discusses ethical use and sets boundaries. Foster open conversations about things like avoiding biased suggestions, not blindly copy-pasting answers without understanding, and maintaining security (e.g., not pasting proprietary code into an unsecured AI service). As AI use grows, new questions will arise (How much can we trust this result? Should AI be writing this critical algorithm?). Make it acceptable to question and critique AI outputs openly. This not only ensures responsibility, but also deepens understanding – developers will dig into why the AI produced a certain output and in doing so, learn more about both the AI and the code. In essence, create a culture where AI is neither feared nor worshipped, but approached with informed enthusiasm and healthy skepticism.

By fostering these practices, you'll nurture a team that's adaptable and ahead of the curve. Companies that successfully integrate AI often have a palpable culture of curiosity and empowerment. Engineers feel encouraged to try new tools, and they aren't afraid to suggest innovative AI-driven approaches to problems. As a leader, you want to unlock that creativity. Remember, the technology will keep evolving – today it's Cursor and Windsurf, tomorrow it might be something even more powerful – so the best long-term investment you can make is in your team's capacity to learn and adapt. That's what an AI-ready engineering culture is all about.

Balancing Automation with Human Judgment

With great power comes great responsibility. As AI takes on a larger role in coding, one of the most important jobs of an engineering manager is to maintain the right balance between automation and human judgment. AI can boost productivity and quality, but unchecked automation can also introduce risks – from subtle bugs to accumulate “tech debt” faster than ever. Here’s how to strike the balance:

- Keep humans in the loop for critical decisions: No matter how advanced the tool, make sure a human engineer is accountable for the final code. AI suggestions should be reviewed and tested just like code written by a junior developer would be. For key architectural decisions or complex algorithm choices, use the AI as a brainstorming partner, but rely on senior engineers to make the final call. The AI might generate five different approaches to a problem in seconds – it’s the team’s job to evaluate which (if any) truly fits the requirements, performance needs, and maintainability standards. Human insight, domain knowledge, and intuition remain invaluable in these moments.
- Establish a code review process that includes AI-generated code: Treat AI contributions as first-class artifacts in your workflow. During code reviews, reviewers should know which parts were AI-generated (developers can indicate this in commit messages or pull request descriptions). This isn’t to stigmatize it, but to ensure reviewers pay extra attention that the code meets the team’s standards and has no hidden flaws. Encourage reviewers to ask “Does this make sense? Do we understand why it’s correct?” just as they would for any code. If the AI wrote a complex bit of logic, the author (developer who applied the suggestion) should be able to explain it – if not, that’s a red flag to go back and refine. This practice ensures automation doesn’t become a black box in your codebase; every line still gets human scrutiny before it hits production.
- Monitor quality metrics for anomalies: As mentioned earlier, keep an eye on indicators like bug rates, code churn, and performance issues after introducing AI. If you notice a spike in quick-fix commits or reverts of AI-generated code, investigate the root cause. It could be that developers are overtrusting the AI and committing code that they later have to rip out. In fact, an analysis by GitClear found that code churn (code discarded shortly after being written) was increasing with heavy AI usage, indicating “AI-induced tech debt” where code is added quickly but then removed soon after. To combat this, you might implement rules like: no AI-generated code goes in without at least one additional pair of eyes, or require tests for any AI-written logic to validate it works as intended. By keeping quality metrics in check, you’ll catch if automation is being overused or misused.

- Use AI to augment, not replace, critical thinking: It's worth reiterating to the team that the goal of these tools is to support their work, not do all their thinking. A good analogy is autopilot in aviation – it can fly the plane in stable conditions, but the pilot is still there to handle turbulence or decide when to change course. Similarly, developers should feel in control: they direct the AI with clear prompts, they verify outputs, and they decide what to accept or reject. If something feels off about an AI suggestion, trust that instinct and investigate further. Sometimes taking a step back and reasoning through the problem (perhaps even without AI) is necessary to validate an approach. Encourage a mindset of collaboration with the AI – it's a colleague with infinite patience and recall, but still learning the ropes. Pair programming paradigms can be useful: treat the AI like a partner whose code you always review.
- Set boundaries for automation: Not everything should be delegated to AI. For example, you might decide that complex domain-specific logic or core infrastructure code must be written and reviewed manually, because the risks of subtle bugs are too high. Or you might limit AI usage in code that requires strict compliance (say, cryptographic algorithms or medical device software) unless the outputs are thoroughly vetted. Communicate these boundaries clearly so developers know when they can lean on AI and when to be more hands-on. Over time, as trust in the tool grows, you can reevaluate and adjust these limits. The key is to start cautious and only expand as confidence builds that automation is handling things well.

In essence, maintain a healthy skepticism and oversight. As one engineering leader put it, "Remember, AI is a tool, not a magic wand". It can dramatically speed up work, but it doesn't eliminate the need for skilled engineers making wise decisions. By embedding that principle in your team's ethos, you ensure that the code produced – whether by human or AI – meets your organization's standards of excellence. The best outcomes arise when automation and human judgment work in harmony, each complementing the other's strengths.

Leading by Example: How Managers and CTOs Can Leverage AI

As a technology leader, your own workflow can benefit from AI tools just as much as your team's can. Moreover, when you use these tools yourself, you set a powerful example that inspires adoption and continuous learning. Here are some ways engineering managers and CTOs can personally leverage AI to stay effective and model the behavior:

- **Use AI for writing and communication:** Engineering leaders often find themselves writing project proposals, design docs, strategy memos, or even HR-style communications. Tools like ChatGPT can be a personal assistant for writing – helping draft documents, refining language, or translating technical jargon into executive-friendly terms. For example, if you need to write a quarterly engineering strategy update, you can prompt the AI with bullet points and ask it to formulate a coherent narrative. It can also help you quickly polish your writing for grammar and tone. By delivering crisp, clear documentation (with the AI's help), you not only save time but also demonstrate to your team that you trust AI assistance in your own work. Just be transparent – there's no shame in mentioning that "I drafted this with a bit of help from an AI tool and then refined it," which can further encourage others to try AI for similar tasks.
- **Automate data analysis and reporting:** Engineering managers often wrangle data from tools like Jira, GitHub, or CI systems to track progress and measure performance. Instead of manually crunching CSV exports in spreadsheets, you can use AI to generate scripts or queries to analyze this data. For instance, if you want to compute DORA metrics or identify bottlenecks in your process, an AI coding assistant can help write a Python script to do it. You provide the prompt ("Connect to Jira API and summarize our team's ticket cycle times for last month"), and the AI gives you a starting script which you can tweak and run. This not only speeds up getting insights, but it also shows your team that you're embracing automation in managerial tasks, not just expecting it from them in coding tasks. Some forward-thinking managers even build their own little AI dashboards or Slack bots – imagine a bot that, every Friday, uses OpenAI's API to compile a brief summary of the week's notable developments or merges. If you experiment with such tools, share them with your peers and team; it signals that generative AI has a place in all aspects of work, including leadership activities.
- **Stay technical and prototype ideas:** Many engineering managers and CTOs want to remain hands-on in small ways to keep their skills sharp and deeply understand the tech their teams use. AI IDEs like Cursor and Windsurf can help you quickly prototype ideas or explore code without consuming too much time. Let's say you have an idea for a new microservice – you can spin up Cursor, describe the high-level idea in the chat, and let the AI scaffold some of the code or setup. In an hour, you might have a rough proof-of-concept

to either discard or pass to a team for further exploration. By doing this, you not only get a tangible feel for what the AI can do, but you also send a message: that experimenting and coding with AI is worthwhile at all levels of seniority. Your team seeing you “get your hands dirty” with AI-driven coding can be motivating (“If my manager can do a quick hack with this, maybe I can use it for my projects too!”). It also gives you credibility when advocating for the team to use the tools – you’ve been in the trenches with it.

- **Use AI for decision support:** Beyond code and writing, AI can assist in problem-solving and decision-making. As a leader, you might face questions like “Should we adopt Technology X or Y?” or “How can we improve our incident response process?”. AI won’t make the decision for you, but you can ask it to outline pros/cons, generate a comparison table, or even draft a strategy based on best practices. For example, you might prompt, “We’re considering moving from monolith to microservices; give me a structured list of potential benefits, risks, and mitigations.” The AI can quickly compile points (sourced from its training on countless tech blogs and articles) that you can then refine with context-specific insights. This accelerates your research phase and ensures you haven’t overlooked common considerations. By using AI as a sounding board, you demonstrate a thoughtful approach to leadership – one that takes advantage of every tool available to make well-informed decisions.
- **Lifelong learning and staying current:** The tech landscape is evolving faster than ever with AI advancements. As a CTO or engineering manager, part of your role is to keep learning so you can guide your organization. Use AI to stay up-to-date: set up alerts or use AI to summarize new research papers, blog posts, or even internal data. Many leaders create a “virtual research assistant” via AI: for instance, feed the executive summary of a new Gartner report into ChatGPT and ask for a brief in plain language. Or if a team publishes an internal post-mortem, use the AI to extract key lessons to share at the next all-hands. Showing that you are leveraging AI to learn and improve sends a strong cultural signal. It says that AI is not just for the devs – it’s something everyone, at every level, should harness for growth and efficiency.

In all these cases, the underlying theme is leading by example. When your team sees you actively using AI tools to be a better manager and technologist, it normalizes the behavior and reduces any stigma or hesitation they might feel. It also equips you with firsthand knowledge to better support your team’s AI usage. You might discover a quirk in Cursor’s agent one day and be able to warn the team, “Hey, watch out for this issue we should account for.” Or you might uncover a brilliant use case for Windsurf that you can propagate to all your developers. By walking the talk, you become not just a manager of AI adoption but a practitioner of it, which is the strongest form of advocacy.

Conclusion

AI is no longer a futuristic add-on to software development – it's here, now, actively changing how we design, code, test, and maintain software. For engineering managers and CTOs, the rise of AI tools presents an enormous opportunity: the chance to supercharge your team's productivity, write more robust code, and create a more enjoyable developer experience. But seizing this opportunity requires thoughtful strategy. It means choosing the right tools (like autonomous agentic IDEs that align with your workflow), introducing them with care and support, and continuously guiding your team to use them effectively and responsibly.

We've explored how AI can transform workflows – making coding faster and catching issues earlier – and how tools like Cursor and Windsurf push the envelope with autonomous, context-aware assistance. The examples and best practices provided illustrate that success isn't just about the technology itself, but about people and process: training your team, fostering an AI-positive culture, measuring impact, and maintaining the balance between AI automation and human expertise. As a leader, you are the catalyst in this transformation. By championing AI adoption, investing in your team's skills, and modeling effective use of these tools, you set the course for your organization to fully leverage what AI has to offer.

In the end, the goal is not to use AI for its own sake, but to build better software and healthier teams. Done right, integrating AI can lead to software being delivered faster with fewer bugs, while engineers find more joy and less toil in their daily work. It's a future where developers focus on creative problem solving and big-picture thinking, supported by their ever-ready AI assistants handling the repetitive grind. As you guide your team into this future, remember that your leadership and vision make all the difference. Embrace the change, stay curious, and lead your engineers in experimenting, learning, and continuously improving. By doing so, you'll not only boost your team's performance – you'll also future-proof your organization in an industry where those who harness new technology intelligently are the ones who leap ahead.

About the Author

[Luciano Ayres](#) is an experienced Engineering Manager and Software Engineer with over two decades of expertise in building scalable, high-impact software solutions. He has led diverse engineering teams, driven innovation in global platforms, and delivered mission-critical systems for some of the world's largest brands, including Nestlé, PepsiCo, and L'Oréal.

As an entrepreneur, Luciano led technical teams in developing mobile applications for international brands, blending engineering excellence with business strategy. He also raised innovation capital to fund and deliver cutting-edge projects, transforming concepts into award-winning software solutions.

Currently, Luciano serves as an Engineering Manager at AB InBev, overseeing backend teams responsible for the world's largest B2B platform, used by 4.5 million monthly users across 28+ countries. He spearheads initiatives to modernize microservices and optimize performance for the Business-to-Business (B2) and Direct-to-Consumer (DTC) platforms.

Luciano holds an MBA in Marketing and a postgraduate degree in Interaction Design. He is also pursuing a postgraduate specialization in Deep Learning, deepening his expertise in AI systems and advanced technologies.

Beyond his technical background, Luciano is passionate about mentorship, emotional intelligence in leadership, and building high-performance engineering cultures. He brings a unique blend of technical excellence and people-centric leadership, which forms the foundation of this book.

When he's not leading engineering teams, Luciano enjoys exploring new technologies, mentoring aspiring developers, and sharing insights on engineering leadership through writing and speaking engagements.

References

1. GitHub Next Survey – Developer report on AI coding tool adoption and benefits. Retrieved from [AI and Its Impact on Engineering Management](#).
2. LinkedIn Pulse: "Level Up Your Dev Team: Introducing AI Code Assistants to the SDLC" – Article by Wagner Souza covering statistics on efficiency and bug reduction, as well as integration tips. Retrieved from [LinkedIn Pulse](#).
3. Cursor Official Blog: "Iterating with Shadow Workspaces" – Discussion on safe autonomous code iteration using shadow workspaces. Retrieved from [Cursor Blog](#).
4. Codeium Windsurf Product Page – Overview of Cascade agent's capabilities including context awareness, command execution, multi-file edits, and autonomous fixes. Retrieved from [Windsurf Editor by Codeium](#).
5. DataCamp Tutorial: "Windsurf AI Agentic Code Editor" – Explanation of Cascade's AutoGPT-like write mode and workflow (approval before run). Retrieved from [DataCamp Tutorial](#).
6. Mabl Engineering Blog: "Empowering Your Engineering Teams for AI Success" – Guidance on training and building a culture of AI literacy. Retrieved from [Mabl Blog](#).
7. DevOps.com: "AI in Software Development: Productivity at the Cost of Code Quality?" – Article highlighting cautionary findings on AI-induced tech debt and code churn. Retrieved from [DevOps.com](#).
8. LeadDev.com: "3 Ways to Use Generative AI for Engineering Management" – Examples of how managers can use AI for writing and automation. Retrieved from [LeadDev.com](#).
9. Cursor Features Page – Highlights of Cursor's agent mode (end-to-end task completion) and other capabilities such as command execution with confirmation. Retrieved from [Cursor Features](#).
10. Windsurf Announcement: "The First Agentic IDE" – Explanation of how the Windsurf Editor combines copilot-style collaboration with independent agent tasks. Retrieved from [Windsurf Announcement](#).